

# Foundations of AI and Machine Learning

Alex Olson

# Welcome!

- My name is Alex Olson
- Senior Research Associate at CARTE
- Bachelor's in AI from the University of Edinburgh
- Master's in AI from UofT in collaboration with the School of Cities
- Published papers in collaboration with a wide array of disciplines
- Work closely with students and faculty on all types of AI

# Wi-Fi Access

- Eduroam Visitor Access (EVA)
- *Permanent* username and password
- Must be re-authorized each day you visit UofT
- New code each day, same phone number

Text **52utoronto**  
To **(833) 338-7626**

# Key AI Technologies

- **Time Series Analysis**
  - Enables systems to predict sequential data
  - Useful for demand forecasting
- **Natural Language Processing (NLP)**
  - Helps machines understand human language
  - Applications in chatbots, sentiment analysis
- **Computer Vision**
  - Allows machines to interpret visual data
  - Used in facial recognition, autonomous vehicles

# AI Maturity Model



## Reactive

Problem-solving focused  
Limited data utilization



## Organized

Centralized data management  
Initial AI projects



## Integrated

AI embedded in multiple business functions  
Advanced analytics capabilities



## Transformative


AI at the core of business strategy  
Continuous innovation and adaptation

# Impact on industries

- Healthcare
  - Drug discovery: Insilico Medicine found new treatments for fibrosis using AI in just 21 days
- Finance
  - Fraud detection: a global bank reduced fraudulent transactions by 50% using AI
- Manufacturing
  - Quality control: Noodle.ai collaborated with a steel mill to deploy an AI application for quality control, reducing suboptimal coil production from 50% to less than 1%

# Why now?

- **Data Availability**
  - Explosion of Big Data
  - Improved Data Storage and Management
- **Computational Power**
  - Advances in GPU Technology
  - Cloud Computing Resources
- **Advanced Algorithms**
  - Breakthroughs in Machine Learning Models
  - Accessibility of Pre-trained Models




# Defining Artificial Intelligence

- Many terms out there with overlapping or confused meanings
  - Artificial Intelligence
  - Machine Learning
  - Deep Learning
  - Data Science
- You will find that in AI, we like to have many terms meaning the same thing!



# Artificial Intelligence

- Getting computers to behave intelligently:
  - Perform **non-trivial tasks** as well as humans do
  - Perform **tasks that even humans struggle with**
- Many sub-goals:
  - Perception
  - Reasoning
  - Control
  - Planning



**My poker face: AI wins multiplayer game for first time**

**Pluribus wins 12-day session of Texas hold'em against some of the world's best human players**



# Speech Recognition: Perception + Reasoning





# Autonomous Driving: Perception + Reasoning Control + Planning





# Game Playing: Reasoning + Planning




# Knowledge-Based AI

Write programs that simulate how people solve the problem

Fundamental limitations:

- Will never get better than a person
- Requires deep domain knowledge
- We don't know how we do some things (e.g., riding a bicycle)



# Data-Based AI = Machine Learning

Write programs that learn the task from examples

- ✓ No need to know how we do it as humans
- ✓ Performance should improve with more examples
- ✗ May need many examples!
- ✗ May not understand how the program works!

# Machine Learning

- Study of algorithms that
  - Improve their performance P
  - At some task T
  - With experience E
- Well defined learning task:  
<P,T,E>

# The Machine Learning Process

- Study of algorithms that
  - Improve their performance P
  - At some task T
  - With experience E
- Well defined learning task:  
<P,T,E>
- Experience
  - Examples of the form  
(input, correct output)
- Task
  - Mapping from input to output
- Performance
  - "Loss function" that measures error w.r.t. desired outcome



# Choices in ML Problem Formulation


- Experience
  - Examples of the form (input, correct output)
- Task
  - Mapping from input to output
- Performance
  - "Loss function" that measures error w.r.t. desired outcome

## Loan Applications

- What historical examples do I have? What is a correct output?
- Predict probability of default? Loan decision? Credit score?
- Do I care more about minimizing False Positives? False negatives?

# What is a “model”?

A useful approximation of the world



Typically, there are **many reasonable models** for the same data

**Training** a model = finding appropriate values for  $(a, b, c, \dots)$

- An **optimization** problem
- “appropriate” = **minimizes the Loss (cost) function**
- We will focus on a common training algorithm later on

# Machine Learning Decisions

- Classification vs Regression
  - Classification: predict between set categories
  - Regression: predict a value (real number)
- Supervised vs Unsupervised
  - Supervised: data with examples of what we want to predict
  - Unsupervised: data but no examples of what we want to predict

# Classification: Three Elements

## 1. Data:

- $x$ : data example with  $d$  attributes
- $y$ : label of example (what you care about)

## 2. Classification model: a function $f_{(a,b,c,\dots)}$

- Maps from  $X$  to  $Y$
- $(a,b,c,\dots)$  are the parameters

## 3. Loss function:

- Penalizes the model's mistakes

Song	Rating
Some nights	★ ★ ★ ★ ★
Skyfall	★
Comfortably numb	★ ★ ★
We are young	★ ★ ★ ★
...	...
...	...
Chopin's 5 <sup>th</sup>	???

# Machine Learning Decisions

- **Classification** vs Regression
  - Classification: predict between set categories
  - Regression: predict a value (real number)
- **Supervised** vs Unsupervised
  - Supervised: data with examples of what we want to predict
  - Unsupervised: data but no examples of what we want to predict

# Classification Loss Function

- How unhappy are you with the answer that the model gave?

- $L_{0-1}(y, f(x)) = \begin{cases} 1 & \text{if: } y \neq f(x) \\ 0 & \text{otherwise} \end{cases}$

- **0-1 loss** function: intuitive but hard to optimize = train



- In practice, we use **approximations** of the 0-1 loss – getting warmer or getting colder

# Regression



## Examples:

- Stock price prediction
- Forecasting epidemics
- Weather prediction



## Regression

What is the temperature going to be tomorrow?



# Machine Learning Decisions

- Classification vs **Regression**
  - Classification: predict between set categories
  - Regression: predict a value (real number)
- **Supervised** vs Unsupervised
  - Supervised: data with examples of what we want to predict
  - Unsupervised: data but no examples of what we want to predict



# Linear Regression

Data:  $S = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_i, y_i)\}$

$\mathbf{x}_i$ : data example with  $d$  attributes


$y_i$ : target of example (what you care about)

**Model:**

$$f(\mathbf{x}; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \dots + \beta_d x_d$$

**Loss function:** Residual Sum of Squares

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \boldsymbol{\beta}))^2$$



# Ridge Regression

- Linear Regression uses all features; model may be complicated
- **Ridge Regression** penalizes large parameter values

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \lambda \sum_{j=0}^d \beta_j^2$$

# Lasso Regression

- As in Ridge Regression, Lasso penalizes large parameters
- Penalizes **absolute** instead of squared coefficient values
- **Zeroes out** more coefficients **BUT** optimization is more involved

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \boldsymbol{\lambda} \sum_{j=0}^d |\beta_j|$$

# Example: Prostate Cancer

Stamey et al. (1989)

- x: cancer volume, prostate weight, age, ...
- y: amount of prostate-specific antigen


Term	LS	Best Subset	Ridge	Lasso
Intercept	2.465	2.477	2.452	2.468
lcavol	0.680	0.740	0.420	0.533
lweight	0.263	0.316	0.238	0.169
age	-0.141		-0.046	
lbph	0.210		0.162	0.002
svi	0.305		0.227	0.094
lcp	-0.288		0.000	
gleason	-0.021		0.040	
pgg45	0.267		0.133	
Test Error	0.521	0.492	0.492	0.479
Std Error	0.179	0.143	0.165	0.164

# Correlation vs. Causation

- **Correlation** measures the strength and direction of a relationship between two variables
- **Causation** refers to a cause-and-effect relationship, where one variable directly influences the other
- It's crucial to remember that a strong correlation doesn't necessarily imply causation


# Correlation vs. Causation

- To avoid confusion between correlation and causation:
  - Consider possible confounding variables or third factors
  - Look for evidence of a causal mechanism
  - Test the relationship using controlled experiments or statistical methods



# Simpson's Paradox

- A trend or relationship between two factors seems to exist when you look at separate groups but disappears or even reverses when you combine the groups together.
- To avoid Simpson's Paradox:
  - Investigate data at different levels of aggregation
  - Consider the influence of confounding variables
  - Use caution when combining data from different sources or groups



# Simpson's Paradox

- In 1973, UC Berkeley found that men applying were more likely to be admitted than women

	All	Men	Women
Applicants Admitted	41%	44%	35%



# Simpson's Paradox


- In 1973, UC Berkeley found that men applying were more likely to be admitted than women
- But when analyzed at a department level, they found only a small subset of departments with a lot of applicants were biased
- Solving the problem required a targeted approach, not a general one

Department	All	Men	Women
A	64%	62%	82%
B	63%	63%	68%
C	35%	37%	34%
D	34%	33%	35%
E	25%	28%	24%
F	6%	6%	7%
Applicants Admitted	39%	45%	30%

# Supervised vs Unsupervised

- So far we have looked at two types of supervised learning
- In both our classification and regression examples, we have examples where we “know the answer”
- With supervised learning, we have a strong definition of the model’s performance


Supervised Learning: Predicted Labels




# Supervised vs Unsupervised

- In unsupervised learning, we don't know what the answer is – collecting this data may be costly, or impossible
- Unsupervised approaches attempt to uncover patterns in the data without relying on a pre-defined label

Unsupervised Learning: Initial Data



Unsupervised Learning: Clustered Data




# Unsupervised Learning: K-Means Clustering

- Clustering approaches seek to uncover groups within data
- Starting with randomly set groups, we measure the similarity of each point to the possible groups, and re-assign
- This process continues until no points change group


# Machine Learning Decisions

- **Classification** vs Regression
  - Classification: predict between set categories
  - Regression: predict a value (real number)
- Supervised vs **Unsupervised**
  - Supervised: data with examples of what we want to predict
  - Unsupervised: data but no examples of what we want to predict


Step 1: Initial Centroids




Step 2: Assign Points to Clusters




Step 2: Assign Points to Clusters




Step 3: Update Centroids



Step 3: Update Centroids



Step 4: Re-Assign Points





# Machine Learning Decisions

- Classification vs Regression
  - Classification: predict between set categories
  - Regression: predict a value (real number)
- Supervised vs Unsupervised
  - Supervised: data with examples of what we want to predict
  - Unsupervised: data but no examples of what we want to predict
- Deep Learning?
  - Methods that use neural networks (this afternoon!)

# Programming Languages for AI



- Python
  - Dominates AI development due to extensive libraries like TensorFlow, PyTorch, and scikit-learn.
  - Easy syntax, strong community, rich ecosystem for AI research and production.
  - Slower execution compared to lower-level languages, can struggle with very large-scale, performance-critical systems.

# Programming Languages for AI



- R
  - Specialized for statistics and data visualization, with packages like caret and ggplot2.
  - Most useful for data preprocessing, statistical modeling, and some ML tasks.
  - Less suited for general-purpose AI development or production-grade systems.

# Programming Languages for AI



- C++
  - High performance for real-time applications like gaming or embedded AI. Libraries like dlib and OpenCV excel in computer vision.
  - Core of many Python-based AI tools (e.g., TensorFlow's backend).
  - Steeper learning curve, verbose syntax, slower development time.
- Java
  - Scalability and enterprise use, with frameworks like Weka and Deeplearning4j.
  - Can interact with Python tools via APIs or frameworks like Apache Spark.
  - Verbose and less favored for rapid prototyping.
- Julia
  - High-performance numerical computing, increasingly adopted for AI and optimization.
  - Growing interoperability with Python (e.g., PyCall).
  - Smaller ecosystem and less community support than Python.

# Machine Learning Libraries PyTorch

- PyTorch
  - Intuitive and flexible, with dynamic computation graphs allowing for easier debugging and experimentation.
  - Strong adoption in research, supported by an active community.
  - Less mature deployment tools compared to TensorFlow (though this gap is narrowing).
  - Can be slower in some production scenarios without optimization.

# Machine Learning Libraries



- TensorFlow
  - Comprehensive ecosystem with tools for training (TensorFlow), deployment (TensorFlow Serving, TensorFlow Lite), and explainability (What-If Tool).
  - TensorFlow.js and TensorFlow Lite make it suitable for web and mobile development.
  - Strong community and corporate support (Google).
  - Integration with Keras offers a high-level API for beginners.
  - Steeper learning curve compared to PyTorch.
  - Debugging can be less straightforward due to static computation graphs (though this has improved with TensorFlow 2.x).

# Machine Learning Libraries



🤗 **Transformers**

- Scikit-learn
  - Easy-to-use interface for classical machine learning tasks like regression, classification, and clustering.
  - Excellent for preprocessing and feature engineering (e.g., PCA, scalers).
  - Strong documentation and wide adoption in education and small-scale projects.
- XGBoost
  - Extremely efficient and scalable for tabular data tasks.
  - Known for achieving high accuracy with minimal tuning.
  - Distributed training support for large datasets.
- HuggingFace Transformers
  - Simplifies the use of pre-trained transformers for NLP, vision, and multimodal tasks.
  - Strong community and regularly updated with state-of-the-art models.
  - Easy fine-tuning and deployment of large language models (LLMs).

# Next: Lab 1

<https://github.com/CARTE-Toronto/mitsubishi-workshop>