

Foundations of AI and Machine Learning

Alex Olson

Welcome!


- My name is Alex Olson
- Senior Research Associate at CARTE
- Bachelor's in AI from the University of Edinburgh
- Master's in AI from UofT in collaboration with the School of Cities
- Published papers in collaboration with a wide array of disciplines
- Work closely with students and faculty on all types of AI

Defining Artificial Intelligence

- Many terms out there with overlapping or confused meanings
 - Artificial Intelligence
 - Machine Learning
 - Deep Learning
 - Data Science
- You will find that in AI, we like to have many terms meaning the same thing!

Artificial Intelligence

- Getting computers to behave intelligently:
 - Perform **non-trivial tasks** as well as humans do
 - Perform **tasks that even humans struggle with**
- Many sub-goals:
 - Perception
 - Reasoning
 - Control
 - Planning



My poker face: AI wins multiplayer game for first time

Pluribus wins 12-day session of Texas hold'em against some of the world's best human players





Speech Recognition: Perception + Reasoning



Autonomous Driving: Perception + Reasoning Control + Planning



Game Playing: Reasoning + Planning




Knowledge-Based AI

Write programs that simulate how people solve the problem

Fundamental limitations:

- Will never get better than a person
- Requires deep domain knowledge
- We don't know how we do some things (e.g., riding a bicycle)



Data-Based AI = Machine Learning

Write programs that learn the task from examples

- ✓ No need to know how we do it as humans
- ✓ Performance should improve with more examples
- ✗ May need many examples!
- ✗ May not understand how the program works!

Machine Learning

- Study of algorithms that
 - Improve their performance P
 - At some task T
 - With experience E
- Well defined learning task: $\langle P, T, E \rangle$

The Machine Learning Process

- Study of algorithms that
 - Improve their performance P
 - At some task T
 - With experience E
- Well defined learning task:
<P,T,E>
- Experience
 - Examples of the form
(input, correct output)
- Task
 - Mapping from input to output
- Performance
 - "Loss function" that measures error w.r.t. desired outcome

Choices in ML Problem Formulation

- Experience
 - Examples of the form (input, correct output)
- Task
 - Mapping from input to output
- Performance
 - "Loss function" that measures error w.r.t. desired outcome

Loan Applications

- What historical examples do I have? What is a correct output?
- Predict probability of default? Loan decision? Credit score?
- Do I care more about minimizing False Positives? False negatives?

Classification: Three Elements

1. Data:

- x : data example with d attributes
- y : label of example (what you care about)

2. Classification model: a function $f_{(a,b,c,\dots)}$

- Maps from X to Y
- (a,b,c,\dots) are the parameters


3. Loss function:

- Penalizes the model's mistakes

Song	Rating
Some nights	★ ★ ★ ★ ★
Skyfall	★
Comfortably numb	★ ★ ★
We are young	★ ★ ★ ★
...	...
...	...
Chopin's 5 th	???

What is a “model”?

A **useful approximation** of the world



Typically, there are **many reasonable models** for the same data

Training a model = finding appropriate values for (a,b,c,...)

- An **optimization** problem
- “appropriate” = **minimizes the Loss (cost)** function
- We will focus on a common training algorithm later on

Classification Loss Function

- How unhappy are you with the answer that the model gave?

- $L_{0-1}(y, f(x)) = \begin{cases} 1 & \text{if: } y \neq f(x) \\ 0 & \text{otherwise} \end{cases}$

- **0-1 loss** function: intuitive but hard to optimize = train

- In practice, we use **approximations** of the 0-1 loss – getting warmer or getting colder



Regression

Examples:

- Stock price prediction
- Forecasting epidemics
- Weather prediction



We will look at:

- Linear Regression
- Ridge Regression
- LASSO



Regression

What is the temperature going to be tomorrow?



Linear Regression

Data: $S = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_i, y_i)\}$

\mathbf{x}_i : data example with d attributes


y_i : target of example (what you care about)

Model:

$$f(\mathbf{x}; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \dots + \beta_d x_d$$

Loss function: Residual Sum of Squares

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \boldsymbol{\beta}))^2$$



Ridge Regression

- Linear Regression uses all features; model may be complicated
- **Ridge Regression** penalizes large parameter values

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \lambda \sum_{j=0}^d \beta_j^2$$

Lasso Regression

- **As in Ridge Regression, Lasso** penalizes large parameters
- Penalizes **absolute** instead of squared coefficient values
- **Zeroes out** more coefficients **BUT** optimization is more involved

Model:

$$f(x; \boldsymbol{\beta}) = \beta_0 + \beta_1 x_1 + \cdots + \beta_d x_d$$

Loss function: Residual Sum of Squares + **penalty** term

$$RSS(\boldsymbol{\beta}) = \sum_{i=1}^n (y_i - f(x_i; \boldsymbol{\beta}))^2 + \lambda \sum_{j=0}^d |\beta_j|$$

Example: Prostate Cancer

Stamey et al. (1989)

- x: cancer volume, prostate weight, age, ...
- y: amount of prostate-specific antigen


Term	LS	Best Subset	Ridge	Lasso
Intercept	2.465	2.477	2.452	2.468
lcavol	0.680	0.740	0.420	0.533
lweight	0.263	0.316	0.238	0.169
age	-0.141		-0.046	
lbph	0.210		0.162	0.002
svi	0.305		0.227	0.094
lcp	-0.288		0.000	
gleason	-0.021		0.040	
pgg45	0.267		0.133	
Test Error	0.521	0.492	0.492	0.479
Std Error	0.179	0.143	0.165	0.164

Correlation vs. Causation

- **Correlation** measures the strength and direction of a relationship between two variables
- **Causation** refers to a cause-and-effect relationship, where one variable directly influences the other
- It's crucial to remember that a strong correlation doesn't necessarily imply causation


Correlation vs. Causation

- To avoid confusion between correlation and causation:
 - Consider possible confounding variables or third factors
 - Look for evidence of a causal mechanism
 - Test the relationship using controlled experiments or statistical methods



Simpson's Paradox

- A trend or relationship between two factors seems to exist when you look at separate groups but disappears or even reverses when you combine the groups together.
- To avoid Simpson's Paradox:
 - Investigate data at different levels of aggregation
 - Consider the influence of confounding variables
 - Use caution when combining data from different sources or groups



Simpson's Paradox

- In 1973, UC Berkeley found that men applying were more likely to be admitted than women

	All	Men	Women
Applicants Admitted	41%	44%	35%

Simpson's Paradox


- In 1973, UC Berkeley found that men applying were more likely to be admitted than women
- But when analyzed at a department level, they found only a small subset of departments with a lot of applicants were biased
- Solving the problem required a targeted approach, not a general one

Department	All	Men	Women
A	64%	62%	82%
B	63%	63%	68%
C	35%	37%	34%
D	34%	33%	35%
E	25%	28%	24%
F	6%	6%	7%
Applicants Admitted	39%	45%	30%

Supervised vs Unsupervised

- So far we have looked at two types of supervised learning
- In both our classification and regression examples, we have examples where we “know the answer”
- With supervised learning, we have a strong definition of the model’s performance


Supervised Learning: Predicted Labels




Supervised vs Unsupervised

- In unsupervised learning, we don't know what the answer is – collecting this data may be costly, or impossible
- Unsupervised approaches attempt to uncover patterns in the data without relying on a pre-defined label

Unsupervised Learning: Initial Data




Unsupervised Learning: Clustered Data




Unsupervised Learning: K-Means Clustering

- Clustering approaches seek to uncover groups within data
- Starting with randomly set groups, we measure the similarity of each point to the possible groups, and re-assign
- This process continues until no points change group


Step 1: Initial Centroids




Step 2: Assign Points to Clusters




Step 2: Assign Points to Clusters




Step 3: Update Centroids



Step 3: Update Centroids



Step 4: Re-Assign Points



Programming Languages for AI



- Python
 - Dominates AI development due to extensive libraries like TensorFlow, PyTorch, and scikit-learn.
 - Easy syntax, strong community, rich ecosystem for AI research and production.
 - Slower execution compared to lower-level languages, can struggle with very large-scale, performance-critical systems.

Programming Languages for AI



- R
 - Specialized for statistics and data visualization, with packages like caret and ggplot2.
 - Most useful for data preprocessing, statistical modeling, and some ML tasks.
 - Less suited for general-purpose AI development or production-grade systems.

Programming Languages for AI



- C++
 - High performance for real-time applications like gaming or embedded AI. Libraries like dlib and OpenCV excel in computer vision.
 - Core of many Python-based AI tools (e.g., TensorFlow's backend).
 - Steeper learning curve, verbose syntax, slower development time.
- Java
 - Scalability and enterprise use, with frameworks like Weka and Deeplearning4j.
 - Can interact with Python tools via APIs or frameworks like Apache Spark.
 - Verbose and less favored for rapid prototyping.
- Julia
 - High-performance numerical computing, increasingly adopted for AI and optimization.
 - Growing interoperability with Python (e.g., PyCall).
 - Smaller ecosystem and less community support than Python.

Machine Learning Libraries



- PyTorch
 - Intuitive and flexible, with dynamic computation graphs allowing for easier debugging and experimentation.
 - Strong adoption in research, supported by an active community.
 - Less mature deployment tools compared to TensorFlow (though this gap is narrowing).
 - Can be slower in some production scenarios without optimization.



TensorFlow

Machine Learning Libraries

- TensorFlow
 - Comprehensive ecosystem with tools for training (TensorFlow), deployment (TensorFlow Serving, TensorFlow Lite), and explainability (What-If Tool).
 - TensorFlow.js and TensorFlow Lite make it suitable for web and mobile development.
 - Strong community and corporate support (Google).
 - Integration with Keras offers a high-level API for beginners.
 - Steeper learning curve compared to PyTorch.
 - Debugging can be less straightforward due to static computation graphs (though this has improved with TensorFlow 2.x).

Machine Learning Libraries



😊 **Transformers**

- Scikit-learn
 - Easy-to-use interface for classical machine learning tasks like regression, classification, and clustering.
 - Excellent for preprocessing and feature engineering (e.g., PCA, scalers).
 - Strong documentation and wide adoption in education and small-scale projects.
- XGBoost
 - Extremely efficient and scalable for tabular data tasks.
 - Known for achieving high accuracy with minimal tuning.
 - Distributed training support for large datasets.
- HuggingFace Transformers
 - Simplifies the use of pre-trained transformers for NLP, vision, and multimodal tasks.
 - Strong community and regularly updated with state-of-the-art models.
 - Easy fine-tuning and deployment of large language models (LLMs).