

Natural Language Process

Alex Olson

Outline

Intro to NLP & Embeddings

Outline

Intro to NLP & Embeddings

Attention, Transformers and GPT

Outline

Intro to NLP & Embeddings

Attention, Transformers and GPT

Applications of LLMs

What is NLP?

- Subfield of AI that focuses on reading, deciphering and producing

What is NLP?

- Subfield of AI that focuses on reading, deciphering and producing human language
- Combines computational linguistics (e.g. rule-based modelling) with statistical, ML, and deep learning approaches

What is NLP?

- Subfield of AI that focuses on reading, deciphering and producing human language
- Combines computational linguistics (e.g. rule-based modelling) with statistical, ML, and deep learning approaches
- Through NLP, machines can understand, analyze and generate text that is meaningful and contextually appropriate

What is NLP?

- Subfield of AI that focuses on reading, deciphering and producing human language
- Combines computational linguistics (e.g. rule-based modelling) with statistical, ML, and deep learning approaches
- Through NLP, machines can understand, analyze and generate human language that is meaningful and contextually appropriate
- While LLMs have driven an explosion in interest, there are many other applications of NLP which paved the way

Common NLP Tasks

- Text Classification, e.g. spam detection
- Named Entity Recognition, e.g. identifying people, places, orga
- Sentiment Analysis, e.g. positive or negative sentiment
- Machine Translation, e.g. Google Translate

Challenges in NLP

- Ambiguity: words can have multiple meanings

Challenges in NLP

- Ambiguity: words can have multiple meanings
- Context: the meaning of a word can change depending on the c

Challenges in NLP

- Ambiguity: words can have multiple meanings
- Context: the meaning of a word can change depending on the c
- Nuance: language is full of subtleties and nuances

Challenges in NLP

- Ambiguity: words can have multiple meanings
- Context: the meaning of a word can change depending on the c
- Nuance: language is full of subtleties and nuances
- Syntax vs Semantics: A phrase can be grammatical but nonsens
ungrammatical but meaningful

Modern NLP

- LLMs have revolutionized NLP

Modern NLP

- LLMs have revolutionized NLP
- Unlike earlier rule-based or statistical models, LLMs learn to generate text through training on large amounts of data

Modern NLP

- LLMs have revolutionized NLP
- Unlike earlier rule-based or statistical models, LLMs learn to generate text through training on large amounts of data
- They can generate text that is coherent and contextually appropriate

Modern NLP

- LLMs have revolutionized NLP
- Unlike earlier rule-based or statistical models, LLMs learn to generate text through training on large amounts of data
- They can generate text that is coherent and contextually appropriate
- They can be fine-tuned for specific tasks

The Encoder-Decoder View of Learning

- Most modern deep learning architectures can be divided into a

The Encoder-Decoder View of Learning

- Most modern deep learning architectures can be divided into a
- In the first step (feature extraction), a model encodes some input representation

The Encoder-Decoder View of Learning

- Most modern deep learning architectures can be divided into a
- In the first step (feature extraction), a model encodes some input into a *latent representation*
- This *latent representation* is meaningful but not directly interpretable

The Encoder-Decoder View of Learning

- Most modern deep learning architectures can be divided into a
- In the first step (feature extraction), a model encodes some input representation
- This *latent representation* is meaningful but not directly interpretable
- In the second step (decoding), the model generates an output based on the latent representation

The Encoder-Decoder View of Learning

- Most modern deep learning architectures can be divided into a
- In the first step (feature extraction), a model encodes some input representation
- This *latent representation* is meaningful but not directly interpretable
- In the second step (decoding), the model generates an output based on the latent representation
- We will first focus on the encoder part of the model

Embeddings

From Real to Symbolic

- Often, machine learning deals with *real-valued* input: data that or can be easily converted to a number and thus contains its own meaning
- Examples: pixel values in an image, audio samples in a sound file, sensor readings
- But what if the input is a symbol?

Symbolic variable

- Text: characters, words, bigrams...
- Recommender Systems: item ids, user ids
- Any categorical descriptor: tags, movie genres, visited resume, product categories...

Symbolic variable

- Text: characters, words, bigrams...
- Recommender Systems: item ids, user ids
- Any categorical descriptor: tags, movie genres, visited resume, product categories...

Symbolic variable

- Text: characters, words, bigrams...
- Recommender Systems: item ids, user ids
- Any categorical descriptor: tags, movie genres, visited resume, product categories...

Symbolic variable

- Text: characters, words, bigrams...
- Recommender Systems: item ids, user ids
- Any categorical descriptor: tags, movie genres, visited resume, product categories...

Notation:

Symbol s in vocabulary V

One-hot representation

$$\text{onehot}(\text{'salad'}) = [0, 0, 1, \dots, 0] \in \{0, 1\}^l$$



One-hot representation

$$\text{onehot}(\text{'salad'}) = [0, 0, 1, \dots, 0] \in \{0, 1\}^{|V|}$$



- Sparse, discrete, large dimension $|V|$
- Each axis has a meaning
- Symbols are equidistant from each other:

$$\text{euclidean distance} = \sqrt{2}$$

Embedding

$$\textit{embedding}(\text{'salad'}) = [3.28, -0.45, \dots 7.1$$

Embedding

$$\textit{embedding}(\text{'salad'}) = [3.28, -0.45, \dots 7.1]$$

- Continuous and dense
- Can represent a huge vocabulary in low dimension, typically:
 $d \in \{16, 32, \dots, 4096\}$
- Axis have no meaning *a priori*
- Embedding metric can capture semantic distance

Embedding

$$\textit{embedding}(\text{'salad'}) = [3.28, -0.45, \dots 7.1]$$

- Continuous and dense
- Can represent a huge vocabulary in low dimension, typically:
 $d \in \{16, 32, \dots, 4096\}$
- Axis have no meaning *a priori*
- Embedding metric can capture semantic distance

Neural Networks compute transformations on continuous vectors

Implementation with Keras

Size of vocabulary $n = |V|$, size of embedding d

```
# input: batch of integers  
Embedding(output_dim=d, input_dim=n, input_length=1)  
# output: batch of float vectors
```

Implementation with Keras

Size of vocabulary $n = |V|$, size of embedding d

```
# input: batch of integers  
Embedding(output_dim=d, input_dim=n, input_length=1)  
# output: batch of float vectors
```

- Equivalent to one-hot encoding multiplied by a weight matrix

$$embedding(x) = onehot(x) \cdot \mathbf{W}$$

Implementation with Keras

Size of vocabulary $n = |V|$, size of embedding d

```
# input: batch of integers  
Embedding(output_dim=d, input_dim=n, input_length=1)  
# output: batch of float vectors
```

- Equivalent to one-hot encoding multiplied by a weight matrix \mathbf{W}

$$embedding(x) = onehot(x) \cdot \mathbf{W}$$

- \mathbf{W} is typically randomly initialized, then tuned by backprop

Implementation with Keras

Size of vocabulary $n = |V|$, size of embedding d

```
# input: batch of integers  
Embedding(output_dim=d, input_dim=n, input_length=1)  
# output: batch of float vectors
```

- Equivalent to one-hot encoding multiplied by a weight matrix

$$embedding(x) = onehot(x) \cdot \mathbf{W}$$

- \mathbf{W} is typically randomly initialized, then tuned by backprop
- \mathbf{W} are trainable parameters of the model

Distance and similarity in Embedding space

Euclidean distance

$$d(x, y) = ||x - y||_2$$

- Simple with good properties
- Dependent on norm (embeddings usually unconstrained)

Distance and similarity in Embedding space

Euclidean distance

$$d(x, y) = \|x - y\|_2$$

- Simple with good properties
- Dependent on norm (embeddings usually unconstrained)

Cosine similarity

$$\text{cosine}(x, y) = \frac{x \cdot y}{\|x\| \cdot \|y\|}$$

- Angle between points in embedding space
- $\text{cosine}(x, y) \in [-1, 1]$
- Expected cosine similarity for random pairs of vectors is 0

Visualizing Embeddings

- Visualizing requires a projection in 2 or 3 dimensions
- Objective: visualize which embedded symbols are similar

Visualizing Embeddings

- Visualizing requires a projection in 2 or 3 dimensions
- Objective: visualize which embedded symbols are similar

PCA

- Limited by linear projection, embeddings usually have complex structure

Visualizing Embeddings

- Visualizing requires a projection in 2 or 3 dimensions
- Objective: visualize which embedded symbols are similar

PCA

- Limited by linear projection, embeddings usually have complex structure

t-SNE

Visualizing data using t-SNE, L van der Maaten, G Hinton, *The Journal of Machine Learning Research*, 2008

t-Distributed Stochastic Neighbourhood Embedding

- Unsupervised, low-dimension, non-linear projection
- Optimized to preserve relative distances between nearest neighbors
- Global layout is not necessarily meaningful


t-Distributed Stochastic Neighbour Embedding

- Unsupervised, low-dimension, non-linear projection
- Optimized to preserve relative distances between nearest neighbours
- Global layout is not necessarily meaningful

t-SNE projection is non deterministic (depends on random initialization)

- Critical parameter: perplexity, usually set to 20, 30
- See <http://distill.pub/2016/misread-tsne/>

Example word vectors



excerpt from work by J. Turian on a model trained by R. Collobert et al